

PP Lab 5

Name: Rajvardhan Reddy Nandyala

Registration Number: 180905093

Section: B

Batch: B1

Roll Number: 19

1.) Write and execute a program in CUDA to add two vectors of length N to meet the following requirements using 3 different kernels.

a) block size as N

b) N threads within a block

c) Keep the number of threads per block as 256 (constant) and vary the number of blocks to handle N elements.

Program:

```
%%cu
#include <cuda.h>
#include <stdlib.h>
#include <stdio.h>

// 1a, 1b, 1c

__global__
void vecAddKernel_1a(float* A, float* B, float* C)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    C[idx] = A[idx] + B[idx];
}

__global__
void vecAddKernel_1b(float* A, float* B, float* C)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
```

```

        C[idx] = A[idx] + B[idx];
    }

__global__
void vecAddKernel_1c(float* A, float* B, float* C,
int n)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    if (idx < n)
    {
        C[idx] = A[idx] + B[idx];
    }
}

```

```

void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float* d_A;
    float* d_B;
    float* d_C;

    cudaMalloc((void**) &d_A, size);
    cudaMalloc((void**) &d_B, size);
    cudaMalloc((void**) &d_C, size);

    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    printf("A: ");
    for (int i = 0; i < n; i++)
    {
        printf("%f, ", A[i]);
    }
    printf("\n");

    printf("B: ");
    for (int i = 0; i < n; i++)
    {
        printf("%f, ", B[i]);
    }
    printf("\n\n");

    vecAddKernel_1a<<<n, 1>>>(d_A, d_B, d_C);
}

```

```

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    printf("A+B (from 1a kernel): ");
    for (int i = 0; i < n ; i++)
    {
        printf("%f, ", C[i]);
    }
    printf("\n");

    vecAddKernel_1b<<<1, n>>>(d_A, d_B, d_C);

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    printf("A+B (from 1b kernel): ");
    for (int i = 0; i < n ; i++)
    {
        printf("%f, ", C[i]);
    }
    printf("\n");

    vecAddKernel_1c<<<ceil(n/256.0), n>>>(d_A, d_B,
d_C, n);

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    printf("A+B (from 1c kernel): ");
    for (int i = 0; i < n ; i++)
    {
        printf("%f, ", C[i]);
    }
    printf("\n");

    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);
}

int main()
{
    float *h_A, *h_B, *h_C;
    int n = 5;

```

```

    int size = n * sizeof(float);

    h_A = (float*) malloc(size);
    h_B = (float*) malloc(size);
    h_C = (float*) malloc(size);

    for (int i = 0; i < n; i++)
    {
        h_A[i] = (i+1) * 10;
        h_B[i] = i+1;
    }

    vecAdd(h_A, h_B, h_C, n);

    return 0;
}

```

Output :

```

A: 10.000000, 20.000000, 30.000000, 40.000000, 50.000000,
B: 1.000000, 2.000000, 3.000000, 4.000000, 5.000000,

A+B (from 1a kernel): 11.000000, 22.000000, 33.000000, 44.000000, 55.000000,
A+B (from 1b kernel): 11.000000, 22.000000, 33.000000, 44.000000, 55.000000,
A+B (from 1c kernel): 11.000000, 22.000000, 33.000000, 44.000000, 55.000000,

```

2.) Write and execute a CUDA program to read an array of N integer values. Sort the array in parallel using parallel selection sort and store the result in another array

Program:

```

%%cu

#include <cuda.h>

#include <stdlib.h>

#include <stdio.h>

__global__

void selectionSortKernel(float* unsorted_arr, float* sorted_arr, int n)

```

```

{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    float key = unsorted_arr[idx];
    int pos = 0;

    for (int i = 0; i < n; i++)
    {
        if (unsorted_arr[i] < key || (unsorted_arr[i] == key && i < idx))
        {
            pos++;
        }
    }

    sorted_arr[pos] = key;
}

```

```

void selectionSort(float* unsorted_arr, float* sorted_arr, int n)
{
    int size = n * sizeof(float);
    float* d_unsorted_arr;
    float* d_sorted_arr;

    cudaMalloc((void**) &d_unsorted_arr, size);
    cudaMalloc((void**) &d_sorted_arr, size);
}

```

```
    cudaMemcpy(d_unsorted_arr, unsorted_arr, size,  
cudaMemcpyHostToDevice);
```

```
selectionSortKernel<<<1, n>>>(d_unsorted_arr, d_sorted_arr, n);
```

```
cudaMemcpy(sorted_arr, d_sorted_arr, size, cudaMemcpyDeviceToHost);
```

```
cudaFree(d_unsorted_arr);
```

```
cudaFree(d_sorted_arr);
```

```
}
```

```
int main()
```

```
{
```

```
    float *h_unsorted_arr, *h_sorted_arr;
```

```
    int n = 5;
```

```
    int size = n * sizeof(float);
```

```
    h_unsorted_arr = (float*) malloc(size);
```

```
    h_sorted_arr = (float*) malloc(size);
```

```
    for (int i = 0; i < 5; i++)
```

```
    {
```

```
        h_unsorted_arr[i] = rand() % 50;
```

```
    }
```

```
    selectionSort(h_unsorted_arr, h_sorted_arr, n);
```

```

printf("unsorted_arr: ");
for (int i = 0; i < n; i++)
{
    printf("%f, ", h_unsorted_arr[i]);
}
printf("\n\n");

printf("sorted_arr: ");
for (int i = 0; i < n; i++)
{
    printf("%f, ", h_sorted_arr[i]);
}
printf("\n");

return 0;
}

```

Output :

```

unsorted_arr: 33.000000, 36.000000, 27.000000, 15.000000, 43.000000,
sorted_arr: 15.000000, 27.000000, 33.000000, 36.000000, 43.000000,

```

3.) Write a execute a CUDA program to read an integer array of size N. Sort this array using odd-even transposition sorting. Use 2 kernels.

Program:

```

%%cu

#include <cuda.h>

```

```

#include <stdlib.h>
#include <stdio.h>

__global__
void oddEven(float* arr, int n)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    if (idx % 2 == 1 && idx + 1 < n)
    {
        if (arr[idx] > arr[idx+1])
        {
            float temp = arr[idx];
            arr[idx] = arr[idx+1];
            arr[idx+1] = temp;
        }
    }
}

```

```

__global__
void evenOdd(float* arr, int n)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    if (idx%2 == 0 && idx+1 < n)
    {
        if (arr[idx] > arr[idx+1])

```



```

    {
        float temp = arr[idx];
        arr[idx] = arr[idx+1];
        arr[idx+1] = temp;
    }
}

```

```

void oddEvenTranspositionSort(float* arr, int n)

```

```

{
    int size = n * sizeof(float);
    float* d_arr;

    cudaMalloc((void**) &d_arr, size);

    cudaMemcpy(d_arr, arr, size, cudaMemcpyHostToDevice);

    for (int i = 0; i <= n/2; i++)
    {
        oddEven<<<1, n>>>(d_arr, n);
        evenOdd<<<1, n>>>(d_arr, n);
    }

    cudaMemcpy(arr, d_arr, size, cudaMemcpyDeviceToHost);

    cudaFree(d_arr);
}

```

```
}  
  
int main()  
{  
    float *h_arr;  
    int n = 5;  
  
    int size = n * sizeof(float);  
  
    h_arr = (float*) malloc(size);  
  
    for (int i = 0; i < 5; i++)  
    {  
        h_arr[i] = rand() % 50;  
    }  
  
    printf("unsorted_arr: ");  
    for (int i = 0; i < n; i++)  
    {  
        printf("%f, ", h_arr[i]);  
    }  
    printf("\n\n");  
  
    oddEvenTranspositionSort(h_arr, n);  
  
    printf("sorted_arr: ");  
    for (int i = 0; i < n; i++)
```

```
{  
    printf("%f, ", h_arr[i]);  
}  
printf("\n");  
return 0;  
}
```

Output :

```
unsorted_arr: 33.000000, 36.000000, 27.000000, 15.000000, 43.000000,  
sorted_arr: 15.000000, 27.000000, 33.000000, 36.000000, 43.000000,
```
