

PP LAB -6 : Programs on Matrix using CUDA

Name : Rajvardhan Reddy Nandyala

Reg No : 180905093

Sem : 6th

Sec : B , Batch : B1

Roll No : 19

1) Write a program in CUDA to add two Matrices for the following specifications:

- a. Each row of the resultant matrix to be computed by one thread.
- b. Each column of the resultant matrix to be computed by one thread.
- c. Each element of the resultant matrix to be computed by one thread.

P1)

```
%%cu
```

```
#include<cuda.h>
```

```
#include<stdio.h>
```

```
#include<stdlib.h>
```

```
__global__ void matAddKernel_1a(float *a,float *b,float *c,int n)
```

```
{
```

```
    int ridA = threadIdx.x;
```

```
    int i;
```

```
    for (i=0;i<n;i++)
```

```
    {
```

```
        c[ridA*n+i] = a[ridA*n+i]+b[ridA*n+i];
```

```
    }
```

```
}
```

```
__global__ void matAddKernel_1b(float *a,float *b,float *c,int m)
```

```
{
```

```
    int col = threadIdx.x;
```

```
    int n = blockDim.x;
```

```
    int i;
```

```
    for (i=0;i<m;i++){
```

```
        c[i*n+col] = a[i*n+col]+b[i*n+col];
```

```
    }
```

```
}
```

```
__global__ void matAddKernel_1c(float *a,float *b,float *c)
```

```
{
```

```
    int row = threadIdx.x;
```

```
    int col = threadIdx.y;
```

```
    int n = blockDim.y;
```

```
    c[row*n+col] = a[row*n+col]+b[row*n+col];
```

```
}
```

```
void matAdd(float *a,float *b,float *c,int m,int n)
```

```
{
```

```
    int size = m*n*sizeof(float);
```

```
    float* d_A;
```

```
    float* d_B;
```

```
    float* d_C;
```

```
    cudaMalloc((void**) &d_A, size);
```

```
cudaMalloc((void**) &d_B, size);
```

```
cudaMalloc((void**) &d_C, size);
```

```
cudaMemcpy(d_A, a, size, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(d_B, b, size, cudaMemcpyHostToDevice);
```

```
int i,j;
```

```
printf("A:\n");
```

```
for (i=0;i<m;i++)
```

```
{
```

```
    for (j=0;j<n;j++)
```

```
    {
```

```
        printf("%f ",*(a + i*n + j));
```

```
    }
```

```
    printf("\n");
```

```
}
```

```
printf("B:\n");
```

```
for (i=0;i<m;i++)
```

```
{
```

```
    for (j=0;j<n;j++)
```

```
    {
```

```
        printf("%f ",*(b + i*n + j));
```

```
    }
```

```
    printf("\n");  
}
```

```
printf("\n");  
matAddKernel_1a<<<1, m>>>(d_A, d_B, d_C, n);  
cudaMemcpy(c, d_C, size, cudaMemcpyDeviceToHost);  
printf("A+B(From the first Kernel):\n");
```

```
for (i=0;i<m;i++)  
{  
    for (j=0;j<n;j++)  
    {  
        printf("%f ",*(c + i*n + j));  
    }  
    printf("\n");  
}
```

```
printf("\n");  
matAddKernel_1b<<<1, n>>>(d_A, d_B, d_C,m);  
cudaMemcpy(c, d_C, size, cudaMemcpyDeviceToHost);
```

```
printf("A+B(From the second Kernel):\n");  
for (i=0;i<m;i++)  
{  
    for (j=0;j<n;j++)  
    {
```

```
        printf("%f ",*(c + i*n + j));  
    }  
    printf("\n");  
}
```

```
printf("\n");  
matAddKernel_1c<<<(1,1), (m,n)>>>(d_A, d_B, d_C);  
cudaMemcpy(c, d_C, size, cudaMemcpyDeviceToHost);
```

```
printf("A+B(From the third Kernel):\n");  
for (i=0;i<m;i++)  
{  
    for (j=0;j<n;j++)  
    {  
        printf("%f ",*(c + i*n + j));  
    }  
    printf("\n");  
}
```

```
printf("\n");
```

```
cudaFree(d_A);  
cudaFree(d_B);  
cudaFree(d_C);  
}
```

```
int main()
{
    float *a,*b,*c;
    int n=3,m=5;
    int size = m*n*sizeof(float);
    a = (float*)malloc(size);
    b = (float*)malloc(size);
    c = (float*)malloc(size);
    int i,j,k=5;
    for (i=0;i<m;i++)
    {
        for (j=0;j<n;j++)
        {
            *(a + i*n + j) = float(k);
            *(b + i*n + j) = float(k+2);
            k+=1;
        }
    }
    matAdd(a,b,c,m,n);
}
```

Output :

```
A:
5.000000 6.000000 7.000000
8.000000 9.000000 10.000000
11.000000 12.000000 13.000000
14.000000 15.000000 16.000000
17.000000 18.000000 19.000000
B:
7.000000 8.000000 9.000000
10.000000 11.000000 12.000000
13.000000 14.000000 15.000000
16.000000 17.000000 18.000000
19.000000 20.000000 21.000000
```

```
A+B(From the first Kernel):
12.000000 14.000000 16.000000
18.000000 20.000000 22.000000
24.000000 26.000000 28.000000
30.000000 32.000000 34.000000
36.000000 38.000000 40.000000
```

```
A+B(From the second Kernel):
12.000000 14.000000 16.000000
18.000000 20.000000 22.000000
24.000000 26.000000 28.000000
30.000000 32.000000 34.000000
36.000000 38.000000 40.000000
```

```
A+B(From the third Kernel):
12.000000 14.000000 16.000000
18.000000 20.000000 22.000000
24.000000 26.000000 28.000000
30.000000 32.000000 34.000000
36.000000 38.000000 40.000000
```

2) Write a program in CUDA to multiply two Matrices for the following specifications:

- Each row of the resultant matrix to be computed by one thread.
- Each column of the resultant matrix to be computed by one thread.
- Each element of the resultant matrix to be computed by one thread.

P2)

```
%%cu
```

```
#include <stdio.h>
```

```

#include <stdlib.h>

#include <cuda.h>

__host__ __device__ void printMatrix(const char* string, int* A, int width)
{
    printf("%s\n", string);
    for (int i = 0; i < width; i++)
    {
        for (int j = 0; j < width; j++)
        {
            printf("%d, ", A[i*width + j]);
        }
        printf("\n");
    }
    printf("\n");
}

__host__ void clearMatrix(int* A, int width)
{
    for (int i = 0; i < width; i++)
    {
        for (int j = 0; j < width; j++)
        {
            A[i*width + j] = 0;
        }
    }
}

__global__

```



```

void multiplyMatrixKernel_2a(int* A, int* B, int* C, int width)
{
    int row = threadIdx.y;
    int k = 0;
    for (int i = 0; i < width; i++)
    {
        k = 0;
        for (int j = 0; j < width; j++)
        {
            k += A[row*width + j] * B[i + width*j];
        }
        C[row*width + i] = k;
    }
}

__global__

```

```

void multiplyMatrixKernel_2b(int* A, int* B, int* C, int width)
{
    int col = threadIdx.x;
    int k = 0;
    for (int i = 0; i < width; i++)
    {
        k = 0;
        for (int j = 0; j < width; j++)
        {
            k += A[i*width + j] * B[col + j*width];
        }
    }
}

```

```

C[i*width + col] = k;
}
}

__global__
void multiplyMatrixKernel_2c(int* A, int* B, int* C, int width)
{
    int col = threadIdx.x;
    int row = threadIdx.y;
    int k = 0;
    for (int i = 0; i < width; i++)
    {
        k += A[row*width + i] * B[col + i*width];
    }
    C[row*width + col] = k;
}

void multiplyMatrix(int* h_A, int* h_B, int* h_C, int width)
{
    int *d_A, *d_B, *d_C;
    int size = width * width * sizeof(int);
    cudaMalloc((void**) &d_A, size);
    cudaMalloc((void**) &d_B, size);
    cudaMalloc((void**) &d_C, size);
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
    dim3 dimBlock(1, 1, 1);

```

```
dim3 dimGrid(1, 1, 1);
dimBlock.x = 1;
dimBlock.y = width;
dimBlock.z = 1;
multiplyMatrixKernel_2a<<<dimGrid, dimBlock>>>(d_A, d_B, d_C, width);
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
printMatrix("A*B: (from 2a kernel): ", h_C, width);
clearMatrix(h_C, width);
cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
dimBlock.x = width;
dimBlock.y = 1;
dimBlock.z = 1;
multiplyMatrixKernel_2b<<<dimGrid, dimBlock>>>(d_A, d_B, d_C, width);
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
printMatrix("A*B: (from 2b kernel): ", h_C, width);
clearMatrix(h_C, width);
cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
dimBlock.x = width;
dimBlock.y = width;
dimBlock.z = 1;
multiplyMatrixKernel_2c<<<dimGrid, dimBlock>>>(d_A, d_B, d_C, width);
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
printMatrix("A*B: (from 2c kernel): ", h_C, width);
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);
```

```
}  
  
int main()  
{  
    int *A, *B, *C;  
    int width = 3;  
    int size = width * width * sizeof(int);  
    A = (int*) calloc(width * width, sizeof(int));  
    B = (int*) calloc(width * width, sizeof(int));  
    C = (int*) calloc(width * width, sizeof(int));  
    int k = 1;  
    for (int i = 0; i < width; i++)  
    {  
        for (int j = 0; j < width; j++)  
        {  
            A[i*width + j] = rand() % 10;  
            B[i*width + j] = rand() % 11;  
            k++;  
        }  
    }  
    printMatrix("A:", A, width);  
    printMatrix("B:", B, width);  
    multiplyMatrix(A, B, C, width);  
    return 0;  
}
```

Output :

```
A:
3, 7, 3,
6, 9, 2,
0, 3, 0,

B:
10, 2, 4,
6, 1, 7,
3, 4, 10,

A*B: (from 2a kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,

A*B: (from 2b kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,

A*B: (from 2c kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,
```